



# Data Mining

## Lab - 2

**Name:** Harmik Rathod

**Enrollment No:** 24010101680

## Numpy & Perform Data Exploration with Pandas

---

### Numpy

1. NumPy (Numerical Python) is a powerful open-source library in Python used for numerical and scientific computing.
2. It provides support for large, multi-dimensional arrays and matrices, along with a collection of mathematical functions to operate on them efficiently.
3. NumPy is highly optimized and written in C, making it much faster than using regular Python lists for numerical operations.
4. It serves as the foundation for many other Python libraries in data science and machine learning, like pandas, TensorFlow, and scikit-learn.
5. With features like broadcasting, vectorization, and integration with C/C++ code, NumPy allows for cleaner and faster code in numerical computations.

### Step 1. Import the Numpy library

```
In [11]: import numpy as np
```

### Step 2. Create a 1D array of numbers

```
In [41]: arr=np.array([10,20,30,40,50])  
arr
```

```
Out[41]: array([10, 20, 30, 40, 50])
```

```
In [15]: arr1=np.arange(11)
arr1
```

```
Out[15]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10])
```

```
In [17]: arr2=np.arange(1,11)
arr2
```

```
Out[17]: array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10])
```

### Step 3. Reshape 1D to 2D Array

```
In [24]: arr3=np.arange(12).reshape(3,4)
arr3
```

```
Out[24]: 2
```

```
In [26]: arr3.dtype #datatype
```

```
Out[26]: dtype('int32')
```

```
In [28]: arr3.size #size of array
```

```
Out[28]: 12
```

```
In [29]: arr3.ndim #dimesional of array
```

```
Out[29]: 2
```

### Step 4. Create a Linspace array

```
In [31]: # np.linspace(start,end,noofele)
arr4=np.linspace(0,5,10)
arr4
```

```
Out[31]: array([0.          , 0.55555556, 1.11111111, 1.66666667, 2.22222222,
                2.77777778, 3.33333333, 3.88888889, 4.44444444, 5.          ])
```

### Step 5. Create a Random Numbered Array

```
In [38]: arr5=np.random.rand(10)
arr5
```

```
Out[38]: array([0.79396546, 0.36140335, 0.86861998, 0.63020115, 0.73472398,
                0.58807095, 0.6907898 , 0.46554542, 0.43849539, 0.46143955])
```

```
In [ ]:
```

## Step 6. Create a Random Integer Array

```
In [44]: arr6=np.random.randint(1,10,size=10)
arr6
```

```
Out[44]: array([5, 2, 5, 1, 6, 9, 1, 2, 7, 1])
```

```
In [ ]:
```

## Step 7. Create a 1D Array and get Max,Min,ArgMax,ArgMin

```
In [47]: arr7=np.random.randint(1,20, size=10)
arr7
```

```
Out[47]: array([19,  2, 14,  3, 15, 13,  2, 11,  5, 12])
```

```
In [48]: arr7.min() #value return
```

```
Out[48]: 2
```

```
In [49]: arr7.max() #value return
```

```
Out[49]: 19
```

```
In [51]: arr7.argmax() #index return
```

```
Out[51]: 0
```

```
In [52]: arr7.argmin() #index return
```

```
Out[52]: 1
```

## Step 8. Indexing in 1D Array

```
In [55]: arr8=np.array([1,2,3,4,5,6,7,8])
```

```
In [56]: arr8[2]
```

```
Out[56]: 3
```

## Step 9. Indexing in 2D Array

```
In [59]: arr9=np.arange(1,13).reshape(3,4)
arr9
```

```
Out[59]: array([[ 1,  2,  3,  4],
                [ 5,  6,  7,  8],
                [ 9, 10, 11, 12]])
```

```
In [61]: arr9[0][0]
```

```
Out[61]: 1
```

```
In [62]: arr9[-1][-1]
```

```
Out[62]: 12
```

```
In [65]: arr9[1][3]
```

```
Out[65]: 8
```

## Step 10. Conditional Selection

```
In [67]: arr10=np.arange(1,17).reshape(4,4)
arr10
```

```
Out[67]: array([[ 1,  2,  3,  4],
                [ 5,  6,  7,  8],
                [ 9, 10, 11, 12],
                [13, 14, 15, 16]])
```

```
In [69]: arr10[arr10>10]
```

```
Out[69]: array([11, 12, 13, 14, 15, 16])
```

```
In [73]: arr10[arr10<10]
```

```
Out[73]: array([1, 2, 3, 4, 5, 6, 7, 8, 9])
```

🔥 You did it! 10 exercises down — you're on fire! 🔥

## Pandas

### Step 1. Import the necessary libraries

```
In [76]: import pandas as pd
```

### Step 2. Import the dataset from this [address](https://raw.githubusercontent.com/justmarkham/DAT8/master/data/users.csv).

```
In [123... users=pd.read_csv('https://raw.githubusercontent.com/justmarkham/DAT8/master/data/u
```

### Step 3. Assign it to a variable called users and use the 'user\_id' as index

```
In [81]: users=pd.read_csv('https://raw.githubusercontent.com/justmarkham/DAT8/master/data/u
users
```

Out[81]:

	age	gender	occupation	zip_code
--	-----	--------	------------	----------

user_id				
1	24	M	technician	85711
2	53	F	other	94043
3	23	M	writer	32067
4	24	M	technician	43537
5	33	F	other	15213
...	...	...	...	...
939	26	F	student	33319
940	32	M	administrator	02215
941	20	M	student	97229
942	48	F	librarian	78209
943	22	M	student	77841

943 rows × 4 columns

## Step 4. See the first 25 entries

In [91]: `users.head(25)`

Out[91]:

	age	gender	occupation	zip_code
user_id				
1	24	M	technician	85711
2	53	F	other	94043
3	23	M	writer	32067
4	24	M	technician	43537
5	33	F	other	15213
6	42	M	executive	98101
7	57	M	administrator	91344
8	36	M	administrator	05201
9	29	M	student	01002
10	53	M	lawyer	90703
11	39	F	other	30329
12	28	F	other	06405
13	47	M	educator	29206
14	45	M	scientist	55106
15	49	F	educator	97301
16	21	M	entertainment	10309
17	30	M	programmer	06355
18	35	F	other	37212
19	40	M	librarian	02138
20	42	F	homemaker	95660
21	26	M	writer	30068
22	25	M	writer	40206
23	30	F	artist	48197
24	21	F	artist	94533
25	39	M	engineer	55107

## Step 5. See the last 10 entries

In [92]: `users.tail(10)`

Out[92]:

	age	gender	occupation	zip_code
<b>user_id</b>				
<b>934</b>	61	M	engineer	22902
<b>935</b>	42	M	doctor	66221
<b>936</b>	24	M	other	32789
<b>937</b>	48	M	educator	98072
<b>938</b>	38	F	technician	55038
<b>939</b>	26	F	student	33319
<b>940</b>	32	M	administrator	02215
<b>941</b>	20	M	student	97229
<b>942</b>	48	F	librarian	78209
<b>943</b>	22	M	student	77841

## Step 6. What is the number of observations in the dataset?

In [87]: `users.shape[0]`

Out[87]: 943

## Step 7. What is the number of columns in the dataset?

In [126...]: `users.shape[1]`

Out[126...]: 1

## Step 8. Print the name of all the columns.

In [88]: `users.columns`

Out[88]: Index(['age', 'gender', 'occupation', 'zip\_code'], dtype='object')

## Step 9. How is the dataset indexed?

In [90]: `users.index`

Out[90]: Index([ 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, ..., 934, 935, 936, 937, 938, 939, 940, 941, 942, 943], dtype='int64', name='user\_id', length=943)

## Step 10. What is the data type of each column?

```
In [94]: users.dtypes
```

```
Out[94]: age          int64
gender        object
occupation    object
zip_code      object
dtype: object
```

## Step 11. Print only the occupation column

```
In [121... users.occupation
```

```
Out[121... user_id
1          technician
2              other
3              writer
4          technician
5              other
...
939         student
940 administrator
941         student
942         librarian
943         student
Name: occupation, Length: 943, dtype: object
```

## Step 12. How many different occupations are in this dataset?

```
In [122... users.occupation.nunique()
```

```
Out[122... 21
```

## Step 13. What is the most frequent occupation?

```
In [102... users.occupation.value_counts().head(1)
```

```
Out[102... occupation
student    196
Name: count, dtype: int64
```

## Step 14. Summarize the DataFrame.

```
In [104... users.describe()
```



Out[104...

	age
<b>count</b>	943.000000
<b>mean</b>	34.051962
<b>std</b>	12.192740
<b>min</b>	7.000000
<b>25%</b>	25.000000
<b>50%</b>	31.000000
<b>75%</b>	43.000000
<b>max</b>	73.000000

## Step 15. Summarize all the columns

In [106...

```
users.describe(include='all')
```

Out[106...

	age	gender	occupation	zip_code
<b>count</b>	943.000000	943	943	943
<b>unique</b>	NaN	2	21	795
<b>top</b>	NaN	M	student	55414
<b>freq</b>	NaN	670	196	9
<b>mean</b>	34.051962	NaN	NaN	NaN
<b>std</b>	12.192740	NaN	NaN	NaN
<b>min</b>	7.000000	NaN	NaN	NaN
<b>25%</b>	25.000000	NaN	NaN	NaN
<b>50%</b>	31.000000	NaN	NaN	NaN
<b>75%</b>	43.000000	NaN	NaN	NaN
<b>max</b>	73.000000	NaN	NaN	NaN

## Step 16. Summarize only the occupation column

In [113...

```
users.age.describe()
```

```
Out[113... count    943.000000
          mean     34.051962
          std      12.192740
          min       7.000000
          25%      25.000000
          50%      31.000000
          75%      43.000000
          max      73.000000
          Name: age, dtype: float64
```

## Step 17. What is the mean age of users?

```
In [112... users.age.mean()
```

```
Out[112... 34.05196182396607
```

## Step 18. What is the age with least occurrence?

```
In [120... users.age.value_counts().tail()
```

```
Out[120... age
7         1
66        1
11         1
10         1
73         1
          Name: count, dtype: int64
```

**You're not just learning, you're mastering it. Keep aiming higher! 🚀**

```
In [ ]:
```